

Reinforcement Learning Lab

Lesson 4: Temporal Difference Methods

Davide Corsi and Alberto Castellini

University of Verona
email: davide.corsi@univr.it

Academic Year 2022-23



UNIVERSITÀ
di VERONA
Dipartimento
di **INFORMATICA**

Environment Setup

The first step for the setup of the laboratory environment is to update the repository and load the **miniconda** environment.

- Update the repository of the lab:

```
cd RL-Lab  
git stash  
git pull  
git stash pop
```

- Activate the *miniconda* environment:

```
conda activate rl-lab
```

Safe Procedure

Always back up the previous lessons' solutions before executing the repository update.

Today Assignment

In today's lesson, we implement the **Q-Learning** and **SARSA** algorithms in Python. In particular, the file to complete is:

`RL-Lab/lessons/lesson_4_code.py`

Inside the file, two functions are partially implemented. The objective of this lesson is to complete them.

- **def QLearning()**
- **def SARSA()**

Expected results can be found in:

`RL-Lab/results/lesson_4_results.txt`

Algorithm: Q-Learning

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Figure: Pseudocode for Q-Learning, from the Sutton and Barto book *Reinforcement Learning: An Introduction*

Algorithm: SARSA

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Figure: Pseudocode for SARSA, from the Sutton and Barto book *Reinforcement Learning: An Introduction*

Assignment Notes

Today's assignment is based on the *DangerousGridWorld* environment, the same as the previous lessons. The suggested assignment's solutions use an exploration function provided in the code, `epsilon_greedy()`. However, any other exploration strategy can be used.

Hint

The class *DangerousGridWorld* (i.e., our environment) comes with a useful function to reset the agent to a random state: `random_initial_state()`. *The suggested assignment's solutions use it as a state-initialization function.*

Results Disclaimer

Given the (high) stochasticity of the method, the obtained results may differ from those suggested. The crucial requirement is to obtain a policy that reaches the goal position.

Pseudocode: Q-Learning

Require: *environment* $[A, S]$, *episodes*, α , γ , *expl_func*, *expl_param*

Ensure: *policy*, *rewards*, *lengths*

```
1:  $\forall a \in A, \forall s \in S$  initialize  $Q(s, a)$  arbitrarily
2:  $\text{rewards}, \text{lengths} \leftarrow [0, \dots, 0]$ 
3: for  $i \leftarrow 0$  to episodes do
4:   Initialize  $s$ 
5:   repeat
6:      $a \leftarrow \text{EXPL\_FUNC}(Q, s, \text{expl\_param})$ 
7:      $s', r \leftarrow \text{take action } a \text{ from state } s$ 
8:      $Q(s, a) \leftarrow Q(s, a) + \alpha(R + \gamma \max_{a' \in A_s} Q(s', a) - Q(s, a))$ 
9:      $s \leftarrow s'$ 
10:   until  $s$  is terminal
11:   Update rewards, lengths
12:  $\pi \leftarrow [0, \dots, 0]$ 
13: for each  $s$  in  $S$  do
14:    $\pi_s \leftarrow \operatorname{argmax}_{a \in A_s} Q(s, a)$ 
15: return  $\pi, \text{rewards}, \text{lengths}$ 
```

▷ Null vectors of length *episodes*

▷ Act and observe
▷ TD

▷ Null vector of length $|S|$
▷ Extract policy

Pseudocode: SARSA

Require: *environment* $[A, S]$, *episodes*, α , γ , *expl_func*, *expl_param*

Ensure: *policy*, *rewards*, *lengths*

```
1:  $\forall a \in A, \forall s \in S$  initialize  $Q(s, a)$  arbitrarily
2:  $rewards, lengths \leftarrow [0, \dots, 0]$ 
3: for  $i \leftarrow 0$  to episodes do
4:   Initialize  $s$ 
5:    $a \leftarrow \text{EXPL\_FUNC}(Q, s, \text{expl\_param})$ 
6:   repeat
7:      $s', r \leftarrow$  take action  $a$  from state  $s$ 
8:      $a' \leftarrow \text{EXPL\_FUNC}(Q, s', \text{expl\_param})$ 
9:      $Q(s, a) \leftarrow Q(s, a) + \alpha(R + \gamma Q(s', a') - Q(s, a))$ 
10:     $s \leftarrow s', a \leftarrow a'$ 
11:   until  $s$  is terminal
12:   Update rewards, lengths
13:  $\pi \leftarrow [0, \dots, 0]$ 
14: for each  $s$  in  $S$  do
15:    $\pi_s \leftarrow \underset{a \in A_s}{\text{argmax}} Q(s, a)$ 
16: return  $\pi, rewards, lengths$ 
```

▷ Null vectors of length *episodes*

▷ Act and observe

▷ TD

▷ Null vector of length $|S|$

▷ Extract policy